

kNN practice

Practicing k-Nearest Neighbors

- Training/Test split
- Try LOOCV to select the best number of neighbors by `knn.cv()`
- Try k-fold CV to select the best number of neighbors by `do.chunk()`
- Compute Training and Test error rates

```
# install.packages("ISLR")
# install.packages("ggplot2") # install.packages("plyr")
# install.packages("dplyr")
# install.packages("class")
# Load libraries
library(ISLR)
library(ggplot2)
library(reshape2)
library(plyr)
library(dplyr)
library(class)
```

Carseats is a simulated data set containing sales of child car seats at 400 different stores. There are 11 variables (3 categorical and 8 numerical).

```
data(Carseats)

str(Carseats)

## 'data.frame':   400 obs. of  11 variables:
##  $ Sales      : num  9.5 11.22 10.06 7.4 4.15 ...
##  $ CompPrice  : num  138 111 113 117 141 124 115 136 132 132 ...
##  $ Income     : num  73 48 35 100 64 113 105 81 110 113 ...
##  $ Advertising: num  11 16 10 4 3 13 0 15 0 0 ...
##  $ Population : num  276 260 269 466 340 501 45 425 108 131 ...
##  $ Price      : num  120 83 80 97 128 72 108 120 124 124 ...
##  $ ShelfLoc   : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2 3 3 ...
##  $ Age        : num  42 65 59 55 38 78 71 67 76 76 ...
##  $ Education  : num  17 10 12 14 13 16 15 10 10 17 ...
##  $ Urban      : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1 ...
##  $ US         : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...

colnames(Carseats)

##  [1] "Sales"      "CompPrice"  "Income"     "Advertising" "Population"
##  [6] "Price"      "ShelveLoc"  "Age"        "Education"   "Urban"
## [11] "US"
```

High is a new binary feature as the response variable where High = No if Sales \leq median(Sales) or High = Yes if Sales $>$ median(Sales). I will drop Sales and the 3 discrete independent variables(ShelveLoc, Urban and US). The goal is to investigate the relationship between High and all the continuous explanatory variables.

```
seats <- Carseats %>% mutate(High=as.factor(ifelse(Sales<=median(Sales), "Low", "High"))) %>% select(-Sales)
str(seats)
```

```
## 'data.frame': 400 obs. of 8 variables:
## $ CompPrice : num 138 111 113 117 141 124 115 136 132 132 ...
## $ Income : num 73 48 35 100 64 113 105 81 110 113 ...
## $ Advertising: num 11 16 10 4 3 13 0 15 0 0 ...
## $ Population : num 276 260 269 466 340 501 45 425 108 131 ...
## $ Price : num 120 83 80 97 128 72 108 120 124 124 ...
## $ Age : num 42 65 59 55 38 78 71 67 76 76 ...
## $ Education : num 17 10 12 14 13 16 15 10 10 17 ...
## $ High : Factor w/ 2 levels "High","Low": 1 1 1 2 2 1 2 1 2 2 ...
```

a. Training/Testing Split

First, I will check the test error where I sample 50% of the observations as a training set, and the other 50% as a test set.

```
# Set random seed
set.seed(333)

# Sample 50% observations as training data
train = sample(1:nrow(seats), 200)
# seats.train = entire rows where the row is specified by the values of train
seats.train = seats[train,]

# The rest 50% as test data
seats.test = seats[-train,]

# YTrain is the observed labels for High on the training set, XTrain is the design matrix
YTrain = seats.train$High
XTrain = seats.train %>% select(-High) %>% scale(center=TRUE, scale=TRUE)

# Test set should be centered and scaled based on the means and variances of the training set
meanvec <- attr(XTrain, 'scaled:center')
sdvec <- attr(XTrain, 'scaled:scale')

# YTest is the observed labels for High on the test set, Xtest is the design matrix
YTest = seats.test$High
XTest = seats.test %>% select(-High) %>% scale(center = meanvec, scale = sdvec)
```

b. Train a kNN classifier and calculate error rates

Now I apply knn() function to train the kNN classifier on the training set and make predictions on training and test sets.

To get the training error, I have to train the kNN classifier on the training set and predict High on the same training set, then I can construct the 2x2 confusion matrix to get the training error rate. I have train=XTrain, test=XTrain, and cl=YTrain in knn(). For now I use k=2.

```
set.seed(444)

# knn - train the classifier and make predictions on the TRAINING set
pred.YTtrain = knn(train=XTrain, test=XTrain, cl=YTrain, k=2)
```

```
# Calculate confusion matrix
conf.train = table(predicted=pred.YTtrain, observed=YTrain)
conf.train
```

```
##           observed
## predicted High Low
##      High    84  10
##      Low     16  90
```

```
# Train accuracy rate
sum(diag(conf.train)/sum(conf.train))
```

```
## [1] 0.87
```

```
# Train error rate
1 - sum(diag(conf.train)/sum(conf.train))
```

```
## [1] 0.13
```

To get the test error, I train kNN on the training set and predict High on the test set, then construct the 2x2 confusion matrix to get the test error rate, where train=XTrain, test=XTest, and cl=YTrain in knn().

```
set.seed(555)
```

```
# knn - train the classifier on TRAINING set and make predictions on TEST set
pred.YTest = knn(train=XTrain, test=XTest, cl=YTrain, k=2)
```

```
# Get confusion matrix
conf.test = table(predicted=pred.YTest, observed=YTest)
conf.test
```

```
##           observed
## predicted High Low
##      High    58  44
##      Low     41  57
```

```
# Test accuracy rate
sum(diag(conf.test)/sum(conf.test))
```

```
## [1] 0.575
```

```
# Test error rate
1 - sum(diag(conf.test)/sum(conf.test))
```

```
## [1] 0.425
```

Test error rate obtained by 2-NN classifier is not good, since 42.5% of the test observations are incorrectly predicted. I should change k, number of neighbors, to improve the error rates. To find optimal value, I will use cross validation.

c. k-fold and Leave-One-Out Cross-validation for selecting best k

- Leave-One-Out Cross-validation (LOOCV)

LOOCV is a special case of k-fold CV in which k is set to equal n. LOOCV has the potential to be expensive to implement, since the model has to be fit n times. This can be very time consuming if n is large, and if each individual model is slow to fit."

```

# validation.error(a vector) to save validation errors in future
validation.error = NULL

# Give possible number of nearest neighbours to be considered
allK = 1:50

# Set random seed to make the results reproducible
set.seed(66)

# For each number in allK, use LOOCV to find a validation error
for (i in allK){ # Loop through different number of neighbors
  pred.Yval = knn.cv(train=XTrain, cl=YTrain, k=i) # Predict on the left out validation set
  validation.error = c(validation.error, mean(pred.Yval!=YTrain)) # Combine all validation errors
}

# Validation error for 1-NN, 2-NN, ..., 50-NN
validation.error

## [1] 0.350 0.375 0.340 0.340 0.325 0.305 0.330 0.310 0.315 0.295 0.290 0.325
## [13] 0.300 0.315 0.295 0.340 0.310 0.320 0.305 0.290 0.320 0.300 0.305 0.320
## [25] 0.310 0.315 0.285 0.275 0.280 0.300 0.295 0.300 0.295 0.310 0.310 0.330
## [37] 0.305 0.320 0.305 0.310 0.300 0.300 0.305 0.300 0.310 0.305 0.295 0.335
## [49] 0.325 0.320

# Best number of neighbors. In the case of a tie, pick larger number of neighbors for simpler model
numneighbor = max(allK[validation.error == min(validation.error)])
numneighbor

## [1] 28

set.seed(67)

# Best k used
pred.YTest = knn(train=XTrain, test=XTest, cl=YTrain, k=numneighbor)

# Confusion matrix
conf.matrix = table(predicted=pred.YTest, true=YTest)
conf.matrix

##           true
## predicted High Low
##      High    73  36
##      Low     26  65

# Test accuracy rate
sum(diag(conf.matrix)/sum(conf.matrix))

## [1] 0.69

# Test error rate
1 - sum(diag(conf.matrix)/sum(conf.matrix))

## [1] 0.31

```

- k-fold Cross-Validation

This approach involves randomly dividing the set of observations into k groups, or folds, of approximately equal size. The first fold is treated as a validation set, and the method is fit on the remaining k - 1 folds.

This procedure is repeated k times; each time, a different group of observations is treated as a validation set and the rest $k - 1$ folds as a training set.

For regression problems, the mean squared error(MSE) is computed on the observations in the k th hold-out fold. This process results in k estimates of the test error: $MSE_1, MSE_2, \dots, MSE_k$. The k -fold CV error is computed by averaging these values,

For classification problems, the error rate is computed on the observations in the k th hold-out fold. Similarly as in regression cases, the CV process yield $Err_1, Err_2, \dots, Err_k$. The k -fold CV error is similarly computed by averaging these values.

Create function `do.chunk()` to select the best number of neighbors using a k -fold CV and to calculate the test error rate afterwards.

```
# do.chunk() returns a data frame consisting of all possible values of folds, each training error and v
do.chunk <- function(chunkid, folddef, Xdat, Ydat, ...){
  train = (folddef!=chunkid) # Get training index
  Xtr = Xdat[train,] # Get training set by the above index
  Ytr = Ydat[train] # Get true labels in training set

  Xvl = Xdat[!train,] # Get validation set
  Yvl = Ydat[!train] # Get true labels in validation set

  predYtr = knn(train=Xtr, test=Xtr, cl=Ytr, ...) # Predict training labels
  predYvl = knn(train=Xtr, test=Xvl, cl=Ytr, ...) # Predict validation labels

  data.frame(fold = chunkid, # k folds
             train.error = mean(predYtr != Ytr), # Training error for each fold
             val.error = mean(predYvl != Yvl)) # Validation error for each fold
}
```

Firstly, I specify $k = 3$ in k -fold CV, and use `cut()` and `sample()` to assign an interval index (1, 2, or 3) to each observation in the training set.

Divide all observations from the training set into 3 intervals, namely, interval 1, 2 and 3. To make the division more random, I sample from all the interval indices without replacement. Call the resulting vector `fold`s.

```
# Specify a 3-fold CV
nfold = 3

# cut: divides all training observations into 3 intervals;
# labels = FALSE instructs R to use integers to code different intervals
set.seed(66)
folds = cut(1:nrow(seats.train), breaks=nfold, labels=FALSE) %>% sample()
folds

## [1] 2 2 3 2 1 3 3 2 2 1 1 2 1 3 1 1 2 2 3 3 2 1 1 2 3 2 3 2 3 1 1 3 2 1 3 2 3
## [38] 3 2 1 1 3 1 3 3 2 3 2 2 1 2 3 3 2 3 1 3 1 1 2 1 2 1 1 2 1 2 3 2 1 2 2 3 3
## [75] 3 3 1 1 1 2 1 3 3 1 1 1 3 3 1 2 1 3 2 3 1 1 2 2 3 3 1 1 3 3 3 2 3 3 2 2
## [112] 3 2 2 3 1 2 2 2 3 3 3 2 1 2 2 2 3 1 2 1 1 1 2 3 3 2 1 2 1 1 1 3 3 2 1 1 2
## [149] 3 1 3 2 3 3 1 1 1 3 1 2 1 2 2 1 3 3 2 3 2 1 1 2 1 2 3 1 1 3 3 2 1 2 2 1 2
## [186] 2 1 2 3 3 2 3 1 1 2 3 1 3 3 1
```

Secondly, use `do.chunk()` to perform a 3-fold CV for selecting the best number of neighbors. The idea is similar to LOOCV, but 3-fold CV instead of doing a 200-fold CV.

```
# Set error.folds (a vector) to save validation errors in future
error.folds = NULL
```

```

# Give possible number of nearest neighbours to be considered
allK = 1:50

set.seed(888)

# Loop through different number of neighbors
for (j in allK){
  tmp = ldply(1:nfold, do.chunk, # Apply do.chunk() function to each fold
              folddef=folds, Xdat=XTrain, Ydat=YTrain, k=j)

  tmp$neighbors = j # Keep track of each value of neighbors

  error.folds = rbind(error.folds, tmp) # combine results
}

```

```

# dim(error.folds)
head(error.folds)

```

```

##   fold train.error val.error neighbors
## 1    1  0.0000000 0.3731343         1
## 2    2  0.0000000 0.3484848         1
## 3    3  0.0000000 0.3134328         1
## 4    1  0.1578947 0.3582090         2
## 5    2  0.2238806 0.3484848         2
## 6    3  0.1654135 0.3731343         2

```

Thirdly, decide the optimal k for kNN based on error.folds.

```

# Transform the format of error.folds for further convenience
errors = melt(error.folds, id.vars=c('fold', 'neighbors'), value.name='error')

# Choose the number of neighbors which minimizes validation error
val.error.means = errors %>%
  # Select all rows of validation errors
  filter(variable=='val.error') %>%
  # Group the selected data frame by neighbors
  group_by(neighbors, variable) %>%
  # Calculate CV error rate for each k
  summarise_each(funs(mean), error) %>%
  # Remove existing group
  ungroup() %>%
  filter(error==min(error))

```

```

# Best number of neighbors
# if there is a tie, pick larger number of neighbors for simpler model
numneighbor = max(val.error.means$neighbors)
numneighbor

```

```
## [1] 31
```

Fourthly, train a 31-NN classifier, and calculate the test error rate.

```

set.seed(99)
pred.YTest = knn(train=XTrain, test=XTest, cl=YTrain, k=numneighbor)

# Confusion matrix
conf.matrix = table(predicted=pred.YTest, true=YTest)

```

```
# Test accuracy rate
sum(diag(conf.matrix)/sum(conf.matrix))

## [1] 0.68

# Test error rate
1 - sum(diag(conf.matrix)/sum(conf.matrix))

## [1] 0.32
```